ASTR4004/8004 Computational Astronomy

INFERENCES

Yuxiang.Qin@anu.edu.au
Woolley Building W22

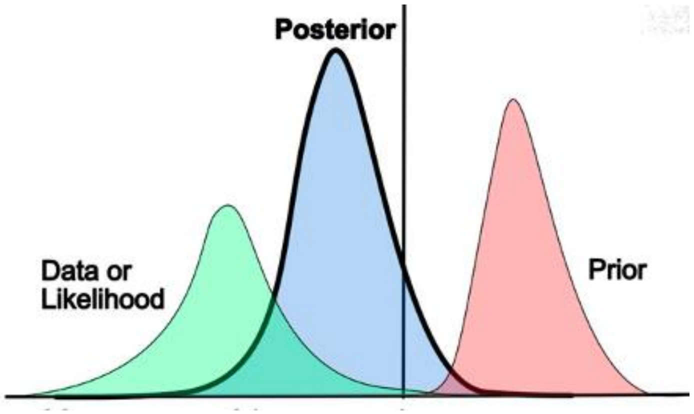

Australian National University

W1-6

- Bash, shell, ssh
- High-Performance Computing
- Python
- C
- Parallel Computing

W7,8,9

- Data Processing
- Statistics
- Plotting
- Regression
- Clustering
- Model Selection

W10

Neural Network

W11

Simulations

W12

Inference

Inferences

using logic and reasoning to draw conclusions or make predictions about unknown information based on available data, evidence, or observations.

Deduction	Induction	Abduction
A conclusion that logically follows from a set of premises, with absolute certainty.	A conclusion based on specific observations, with some degree of uncertainty.	A conclusion based on the best explanation for a set of observations, with some degree of uncertainty.
All men are mortal. Socrates is a man. Therefore, Socrates is mortal. – Aristotle	A turkey found that every day for the past year he was fed at 9am sharp. Therefore, he concludes that "I am always fed at 9am". However, on the morning of Christmas eve, he had his throat cut. – Bertrand Russell	I had grasped the significance of the silence of the dog, for one true inference invariably suggests others.... Obviously, the midnight visitor was someone whom the dog knew well. – <The Adventure of Silver Blaze>
If a star's mass is above 10 solar masses, it will end its life in a supernova explosion. R136a1 has a mass of 200 solar masses. Therefore, R136a1 will end its life in a supernova explosion.	Many observed exoplanets have been found to orbit close to their host stars. It's likely that the newly discovered exoplanet Kepler-452b also orbits close to its star.	We observe the circular velocity of stars and gas in many galaxies remains flat in the outskirt of galaxies. Dark matter!

Bayesian Inference

updating our understanding based on new evidence

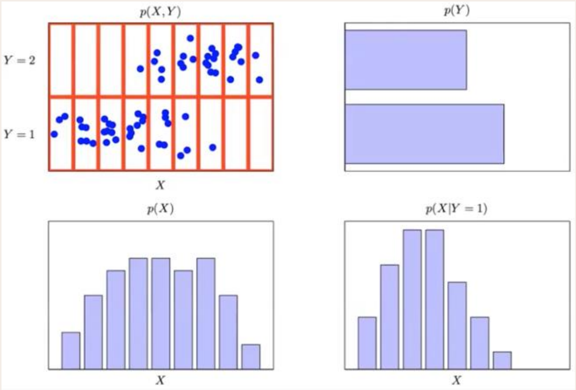

- Joint probability $\mathrm{P(X, Y)}$: the probability of X and Y

Recap W7

- Marginal probability $\mathrm{P(X)}$ or $\mathrm{P(Y)}$: the probability of X or Y
- Conditional probability $\mathrm{P(X|Y)}$ or $\mathrm{P(Y|X)}$: the probability of X given Y or of Y given X
- Sum rule, $\mathrm{P(X)} = \sum_Y \mathrm{P(X, Y)}$, $\mathrm{P(Y)} = \sum_X \mathrm{P(X, Y)}$
- Product rule, $\mathrm{P(X, Y)} = \mathrm{P(X|Y)P(Y)} = \mathrm{P(Y|X)P(X)}$: the probability of X and Y equals to the probability of X (or Y) times the probability of Y given X (or of X given Y)
- Bayes' rule: $\mathrm{P(X|Y)} = \mathrm{P(Y|X)P(X)/P(Y)}$

X: the person having COVID; Y: the person tested positive on a RAT
Say 1% of people in this population are expected to have COVID, 90% of people with COVID will test positive on a RAT, 5% of all people tested positive (regardless of whether they have COVID or not)

i.e., P(having COVID) = 0.01, P(tested positive | having COVID) = 0.9, P(tested positive) = 0.05
P(having COVID | tested positive) 0.01 * 0.9 / 0.05 = 0.18

The Monty Hall problem

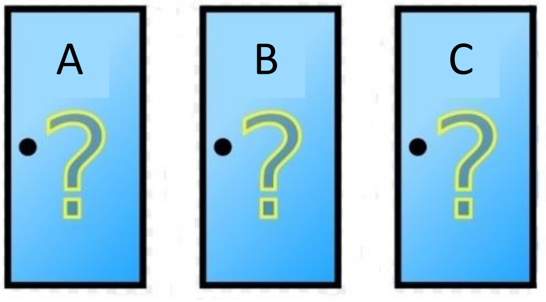

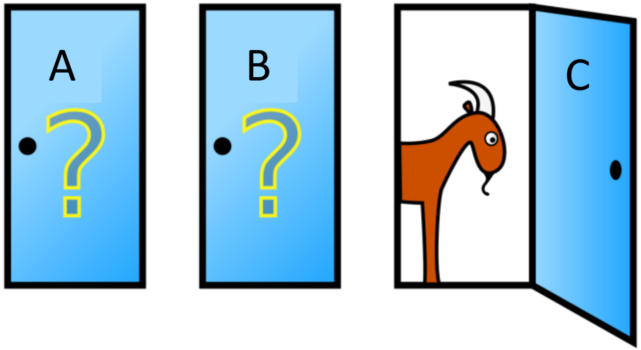

$$P(A|C) = P(C|A)P(A)/P(C)$$

P(A): the probability of the prize behind Door A (1/3) -> Prior

P(C|A): If the prize is behind Door A, the probability of the host opening Door C is 0.5

P(C): the probability of the host opening Door B:

- P(C |A) = 0.5
- P(C | not A):
  - P(C | If the prize is behind door B) = 1
  - P(C | If the prize is behind door C) = 0.

P(C) = (0.5+1+0) / 3 = 0.5

P(A|C) = P(C|A)P(A)/P(C) = 1/3 -> Posterior

Bayesian Inference

updating our understanding based on new evidence

X: your model parameters θ
Y: the observed data D

Say you look at the data points $(x_1, x_2, \dots, x_N)$ and find that they seem to follow a normal distribution: $\theta \equiv (\mu, \sigma)$

$$\mathrm{P}(D|\theta) = \prod_{i=1}^{N} \frac{1}{\sqrt{2\pi}\sigma} e^{-\left(\frac{x_i-\mu}{\sqrt{2}\sigma}\right)^2}$$

Bayes' rule: $\mathrm{P}(\theta|D) = \mathrm{P}(D|\theta)\mathrm{P}(\theta)/\mathrm{P}(D)$

Recap W7

- Joint probability $\mathrm{P}(\mathrm{X}, \mathrm{Y})$: the probability of X and Y
- Marginal probability $\mathrm{P}(\mathrm{X})$ or $\mathrm{P}(\mathrm{Y})$: the probability of X or Y
- Conditional probability $\mathrm{P}(\mathrm{X}|\mathrm{Y})$ or $\mathrm{P}(\mathrm{Y}|\mathrm{X})$: the probability of X given Y or of Y given X
- Sum rule, $\mathrm{P}(\mathrm{X}) = \sum_Y \mathrm{P}(\mathrm{X}, \mathrm{Y})$, $\mathrm{P}(\mathrm{Y}) = \sum_X \mathrm{P}(\mathrm{X}, \mathrm{Y})$
- Product rule, $\mathrm{P}(\mathrm{X}, \mathrm{Y}) = \mathrm{P}(\mathrm{X}|\mathrm{Y})\mathrm{P}(\mathrm{Y}) = \mathrm{P}(\mathrm{Y}|\mathrm{X})\mathrm{P}(\mathrm{X})$: the probability of X and Y equals to the probability of X (or Y) times the probability of Y given X (or of X given Y)
- Bayes' rule: $\mathrm{P}(\mathrm{X}|\mathrm{Y}) = \mathrm{P}(\mathrm{Y}|\mathrm{X})\mathrm{P}(\mathrm{X})/\mathrm{P}(\mathrm{Y})$

- $\mathrm{P}(D|\theta)$: the likelihood function, representing the probability of observing the data (D) given a specific set of modeling parameters (θ)
- $\mathrm{P}(\theta)$: the prior probability distribution over the modeling parameters (θ). This represents our initial beliefs of possible values of θ
- $\mathrm{P}(D) \equiv \sum_\theta \mathrm{P}(D, \theta) = \sum_\theta \mathrm{P}(\theta)\mathrm{P}(D|\theta) = \int d\theta\, \mathrm{P}(\theta)\mathrm{P}(D|\theta)$: the marginal likelihood, representing the overall probability of observing the data (D)
- $\mathrm{P}(\theta|D)$: the posterior probability distribution over the modeling parameters (θ) given the data.

Bayesianism vs Frequentism

To a frequentist, probability does not represent degree of our belief.
Instead, it is only the frequency of a long-run experiment.

Example 1: When a frequentist flips a coin many many (infinitely many) times, half will be heads. But when a Bayesian statistician flips a coin many times and obtain half is head, they believe the probability for the next flip being head is half.

Example 2: When a frequentist is asked to determine the probability of flipping a coin to land on the head, they will keep flipping the coin many many times, collect data and calculate the frequency of head in the end of the experiment. But a Bayesian statistician will first guess e.g., 50% based on their prior knowledge of coin flipping and keep update the probability based on what they observe when flipping the coin

Example 3: When I flip one coin and hide the result, a Bayesian statistician will still say based on my prior knowledge, the probability of the coin being head is 50% and a frequentist will refuse to answer the question and say there is no probability because there is no long-run frequency, i.e., a probability.

Example 4: Pooling experts say there is a 51% probability this candidate will win the election while a frequentist will yell at them and say it makes no sense because there is only one election, so we cannot define a long-run frequency.

Monty Hall meets frequentists

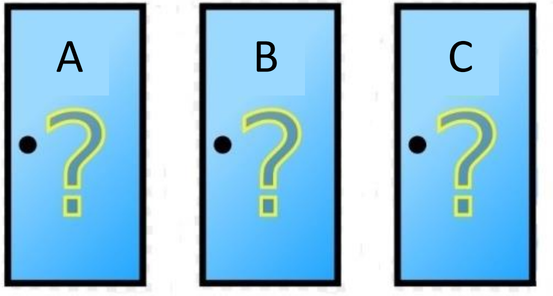


```
def host_open_door(my_choice, win_door, door_options):
    # Determine which door the host opens
    for door in door_options:
        # cannot open my choice or the winning door
        if door != my_choice and door != win_door:
            return door


def my_switches(host_door, my_choice, door_options):
    # Simulate the contestant switching doors
    for door in door_options:
        # cannot be the original choice or the host door
        if door != my_choice and door != host_door :
            return door
```

```
door_options = ["A", "B", "C"]
win_by_switching = []
win_by_staying = []

# Simulate many games of Monty Hall
for game_num in range(10000):
    # Randomly choose the winning door
    win_door = random_choice(door_options)

    # Randomly choose a door for the contestant to initially choose
    my_choice = random_choice(door_options)

    # Host opens one of the other two doors that is not the winning one
    host_door = host_open_door(my_choice, win_door, door_options)

    # contestant switches doors if they want
    switch_door = switches(host_door, my_choice, door_options)

    if switch_door == win_door:
        win_by_switching.append(1)  # Switching won this game
        win_by_staying.append(0)  # Staying lost this game
    else:
        win_by_switching.append(0)  # Switching lost this game
        win_by_staying.append(1)  # Staying won this game

# Calculate the probabilities of winning
prob_win_by_switching = sum(win_by_switching)/10000
prob_win_by_staying   = sum(win_by_staying)/10000
```

Bayesian inference pseudocode

Say you look at the data points $(x_1, x_2, \ldots x_N)$ and find that they seem to follow a normal distribution: $\theta \equiv (\mu = 5, \sigma)$

```
# Initial prior and samples
prior = initial_prior
std_samples = sample_parameters(initial_prior, std_min, std_max, num_samples)
mean_guessed = 5

prior_samples      = np.zeros(num_samples)
likelihood_samples  = np.zeros(num_samples)
posterior_samples   = np.zeros(num_samples)

# Inference loop
for istep in range(num_steps):

  for i in range(num_samples):
    prior_samples[i] = prior(std_samples[i])
    likelihood_samples[i] = np.prod(norm_pdf(data, mean_guessed, std_samples[i]))
    posterior_samples[i] = likelihood_samples[i] * prior_samples[i]

  marginal_likelihoods = np.sum(likelihood_samples) / num_samples
  posterior_samples   /= marginal_likelihoods

  # Interpolate posterior for next step's prior
  f_interp = interp1d(std_samples, posterior_samples)
  def prior(mean, std):
    return f_interp(mean , std)

  std_samples = sample_parameters(prior, std_min, std_max, num_samples)
```

$\mathrm{P}(\theta|D) = \mathrm{P}(D|\theta)\mathrm{P}(\theta)/\mathrm{P}(D)$

$\mathrm{P}(D|\theta)$: likelihood

$\mathrm{P}(\theta)$: prior probability

$\mathrm{P}(D)$: marginal likelihood

$\mathrm{P}(\theta|D)$: posterior probability

Points in a sample might only have a uniform volume when initialized to be uniform and are very likely not the case at later times!

open inference-norm1d.ipynb

Bayesian evidence

$\mathrm{P}(D) \equiv \sum_\theta \mathrm{P}(D,\theta) = \sum_\theta \mathrm{P}(\theta)\mathrm{P}(D|\theta) = \int d\theta\, \mathrm{P}(\theta)\mathrm{P}(D|\theta)$:

The marginal likelihood, representing the overall probability of observing the data (D)

$$P(D|\mathcal{M}_1) = \int d\theta\, \mathrm{P}(\theta|\mathcal{M}_1)\mathrm{P}(D|\theta,\mathcal{M}_1) \qquad P(D|\mathcal{M}_2) = \int d\theta\, \mathrm{P}(\theta|\mathcal{M}_2)\mathrm{P}(D|\theta,\mathcal{M}_2)$$

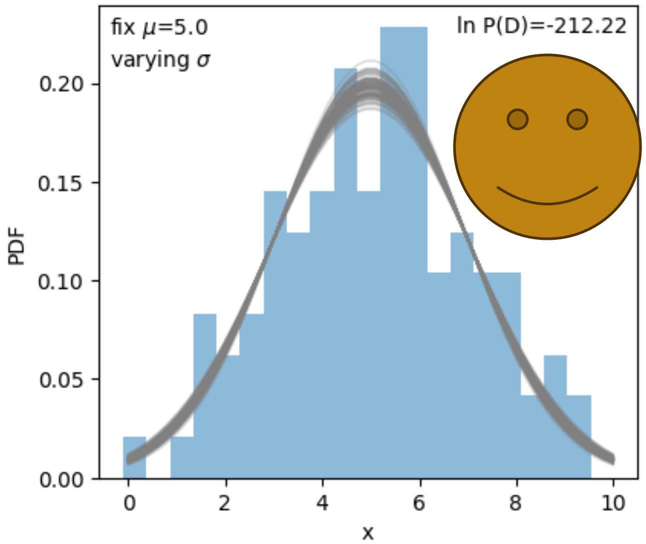

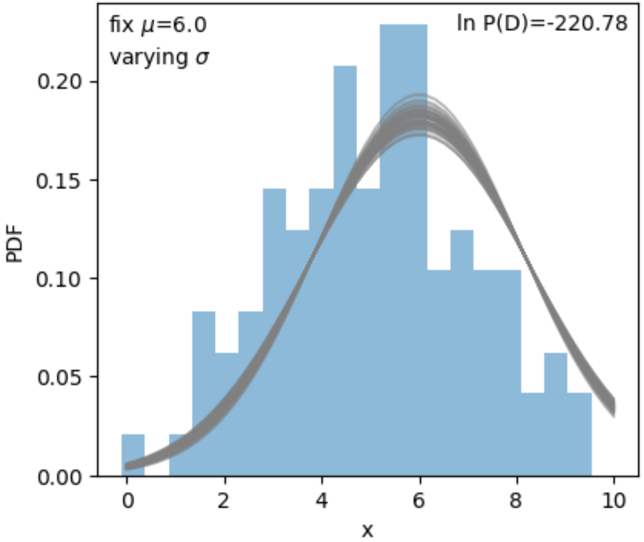

A higher evidence indicates the model is more likely to have generated the observed data.

open inference-norm2d.ipynb

prerequisite

We will start at 1:05pm, before that please

git clone/pull from https://github.com/qyx268/astr4004/

Install matplotlib, numpy, astropy, hmf, scipy, commah,
Install emcee, dynasty, corner